

Large Language Models as Hyperparameter Optimizers

Stephanie Zhou, Jeena Mahajan, Dredheza Gashi Baraliu

MATH 467 – Professor Alcala

Motivation: LLMs & Hyperparameter Tuning

- Hyperparameter tuning: critical for optimization performance
 - Bad parameters = bad performance
 - Wasted iterations
 - Divergence
- Traditional search methods like grid search and random search
 - Costly
 - Can be inefficient
- LLMs: have learned patterns across massive amounts of data
- **Opportunity:** Use LLMs to suggest hyperparameter values for optimization problems


Research Question: Can LLM's suggest effective hyperparameters for optimization algorithms, outperforming traditional search methods?

Gradient Descent

- Purpose: the optimization algorithm we'll use to minimize the benchmark function
- Core Idea:
 - Goal: minimize the function $f(x)$
 - Gradient: direction of steepest increase
 - To minimize, move in the opposite direction
- Iterative update rule: move a small amount “ α ” downhill

$$x_{\{k+1\}} = x_k - \alpha \nabla f(x_k)$$

- Intuition for Gradient Descent = repeatedly taking small downhill steps until you reach the “valley”





Gradient Descent: Hyperparameters

$$v_{t+1} = \beta v_t - \alpha \nabla f(x_t)$$

$$x_{t+1} = x_t + v_{t+1}$$

- Learning rate (α):
 - Too small → very slow
 - Too large → overshoot or diverge
 - “Just right” → fast and stable convergence
 - Intuition: how “big” of a step you take downhill
- Momentum: vector v
 - Adds memory of past steps to smooth the path.
 - v “remembers” past gradients
 - Helps push through small bumps or noisy slopes.
 - Prevents getting stuck in shallow valleys.




Benchmark Functions

- Rastrigin Function

- $f(x, y) = 20 + (x^2 - 10 \cos 2\pi x) + (y^2 - 10 \cos 2\pi y)$, global min at $(0, 0)$, $f = 0$
- Highly multi-modal, many shallow local minima
=> easy to get trapped
- In our experiments:
GD in \mathbb{R}^2 , $x_0 = (-1.5, 2)$
- Used to test whether our learning-rate choices keep gradient descent moving smoothly on a bumpy surface


Rastrigin Function – 3D Surface Plot



Benchmark Functions

- Rosenbrock Function
 - $f(x, y) = (1 - x)^2 + 100(y - x^2)^2$, global min at $(1, 1)$, $f = 0$
 - Narrow, curved valley => step size-sensitive
 - In our experiments:
 - GD in \mathbb{R}^2 , $x_0 = (-1.5, 2)$ [randomly sampled]
 - Used to test whether hyperparameters allow fast, stable convergence along the valley

Rosenbrock Function Surface



Methods

1. Grid Search

- Test preset (learning rate, momentum) pairs
- Run gradient descent with momentum; keep the lowest final value
- LR: [1e-4, 1e-3, 1e-2, 1e-1]; Momentum: [0.0, 0.5, 0.9]

2. Random Search

- Sample 4 random (lr, momentum) pairs
- Run gradient descent for each; pick the best performer
- LR: [6.2e-04 → 1.6e-03]; Momentum: [0.431 → 0.539]

3. LLM-Guided Optimization




- LLM receives function info + valid lr/m ranges
- Returns 3 candidate (lr, momentum) pairs
- Run gradient descent for each; choose the lowest final value
- LR: [1.0e-04 → 1.0e-03]; Momentum: [0.445 → 0.850]

AI Models Used

1. OpenAI GPT-4o
 - Higher-capability model
2. OpenAI GPT-4o-mini
 - Lighter, faster, cheaper variant


Results: Grid Search (Rosenbrock)

- “Best” combination was lr =0.001, momentum = 0.9
 - Best: lowest function value (2.387590e-09) in smallest number of iterations: **2020**
- Observations:
 - Very small learning rate: slow convergence, takes up all 5000 iterations & doesn’t reach lowest value
 - High learning rates: diverge
 - Momentum + Average Learning Rate:
 - Builds up velocity across the curved valley, but learning rate is small enough to not overshoot the valley
- Performed a “refined” grid search around this optimal to see if we could achieve a better optimal hyperparameter combination
 - Result: Learning Rate: 0.0014, Momentum: 0.93; Iterations: **925**



Results: Grid Search (Rastrigin)

- “Best” combination was lr =0.001, momentum = 0
 - Best: reaches local minimum (4.974790) in smallest number of iterations: **36**
- Observations:
 - Very small learning rate: slow convergence, takes up all 5000 iterations & doesn’t reach lowest value
 - High learning rates: diverge -> avoid the many “small basins” due to large jumps
 - Momentum: hurts in this case
 - Without momentum, GD moves cautiously and avoids oscillations in the rugged landscape
 - With momentum, GD overshoots small basins and gets stuck for longer.
- Performed a “refined” grid search around this optimal to see if we could achieve a better optimal hyperparameter combination
 - Result: Learning Rate: 0.0014, Momentum: 0.03; Iterations: **24**



Results: Random Search (Rosenbrock)

- Performance Summary
 - Highly unstable across seeds
 - Only seed 0 produced a near-optimal value
 - Seeds 1 and 2 failed to converge well
 - All runs hit the full 2000 iterations

Seed	Best LR	Momentum	Final Value	Iterations
0	3.6001e-04	0.861	0.001826	2000
1	2.5639e-05	0.869	0.815882	2000
2	5.1045e-04	0.489	0.021873	2000


Results: Random Search (Rastrigin)

- Performance Summary
 - Very unstable across sampled hyperparameters
 - Fast convergence: ~44-51 iterations
 - Seed 0 performs the best: reaches global minimum

Seed	Best LR	Momentum	Final Value	Iterations
0	1.6052e-03	0.539	1.6052e-03	51
1	6.2163e-04	0.533	0.994959	51
2	1.2570e-03	0.431	3.97983	44

Results: LLMs (Rosenbrock)

- Performance Summary
 - GPT-4o & GPT-4o-mini are extremely stable
 - LR: 5×10^{-4}
 - Momentum = 0.80
 - Final values consistently between 0.00141-0.00227
 - All runs use the full budget (slow convergence)




Results: LLMs (Rosenbrock)

Seed	Model	LR	Momentum	Final Value	Iterations
0	GPT-4o	5.0e-04	0.800	0.002119	2000
0	GPT-4o-mini	5.0e-04	0.800	0.002119	2000
1	GPT-4o	5.0e-04	0.800	0.002272	2000
1	GPT-4o-mini	5.0e-04	0.800	0.002272	2000
2	GPT-4o	5.0e-04	0.800	0.001407	2000
2	GPT-4o-mini	5.0e-04	0.800	0.001407	2000

Results: LLMs (Rastrigin)

- Performance Summary
 - Seeds 0 and 1 reached the same final value
~0.994959 (~100 iterations)
 - Consistent across seeds 0 and 1, but differs for seed 2
- GPT-4o:
LR = 5e-4 and 8e-4
Momentum = 0.80 and 0.85.
- GPT-4o-mini: LR = 1e-4, Momentum = 0.70



Results: LLMs (Rastrigin)

Seed	Model	LR	Momentum	Final Value	Iterations
0	GPT-4o	8.0e-04	0.850	0.994959	151
0	GPT-4o-mini	1.0e-04	0.700	0.994959	98
1	GPT-4o	5.0e-04	0.800	0.994959	101
1	GPT-4o-mini	1.0e-04	0.700	0.994959	101
2	GPT-4o	8.0e-04	0.850	3.97983	151
2	GPT-4o-mini	1.0e-04	0.700	3.97983	100

Future Experimentation

- Extend experiments to additional benchmark functions
 - Test on functions with different geometries
 - Verifies whether the LLM generalizes beyond the current benchmark functions
- Evaluate alternative LLMs and use prompt engineering
 - Compare different models
 - Vary prompt structure, details and constraints to optimize the prompt itself
 - Determine whether performance reliant on the prompt

Conclusion

- **Rosenbrock:** Grid Search did the best; **Rastrigin:** Random Search did the best
 - Random Search is stochastic: might miss the “sweet” spot for the valley for Rosenbrock, but can “luck” into hyperparameters that escape local minima for Rastrigin
- LLMs produced reasonable, conservative learning rates & momentum values
 - Can be promising in higher-dimensional spaces where we need guidance without expensive computation
- On these benchmarks, simple baselines were already very strong
- GPT-4o-mini was often the most competitive LLM
- This sets up a foundation for more complex, realistic experiments



Thank you!